Diffusion Limited Aggregation

Candidate Number: [Your Number]

Hilary Term 2026

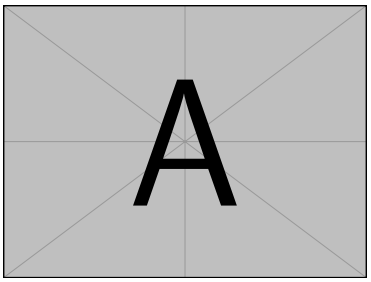

A dissertation submitted for the degree of
Mathematics Part B

Abstract

add abstract here (or delete)

Contents

1 Introduction **5**

1.1 Overview . . . 5

1.2 Theoretical Context . . . 5

1.3 Fractal Dimension . . . 6

1.4 Anisotropy and Thesis Goals . . . 6

2 Potential Theory **7**

2.1 Preliminaries . . . 7

2.1.1 Martingales and the Optional Stopping Theorem . . . 7

2.1.2 Measure Theory Concepts . . . 8

2.1.3 Conformal Maps . . . 9

2.2 Discrete Potential Theory . . . 9

2.2.1 Simple Random Walks and the Generator . . . 9

2.2.2 The Discrete Dirichlet Problem . . . 10

2.2.3 Recurrence and the Potential Kernel . . . 12

2.2.4 Discrete Harmonic Measure . . . 13

2.3 Continuous Potential Theory . . . 15

2.3.1 Brownian Motion and the Laplacian . . . 15

2.3.2 Harmonic Measure . . . 16

2.3.3 Conformal Invariance . . . 17

2.3.4 Solvable Geometries . . . 18

3 Algorithms **21**

3.1 Introduction . . . 21

3.2 Initialisation . . . 21

3.2.1 Off-Lattice: The Launching Circle . . . 21

3.2.2 On-Lattice: The Fuzzy Annulus . . . 22

3.3 Far-Field Diffusion . . . 22

3.3.1 Off-Lattice: Re-injection from Infinity . . . 22

3.3.2 On-Lattice: Walk-to-Line . . . 23

3.4 Near-Field Diffusion . . . 24

3.4.1 Off-Lattice: Walk-to-Circle . . . 24

3.4.2 On-Lattice: Walk-to-Square . . . 25

3.5 Aggregation . . . 25

3.5.1 On-Lattice: Sticky Boundary . . . 25
3.5.2 Off-Lattice: Exact Adhering . . . 25
3.6 The Hybrid Model . . . 27
3.6.1 Motivation and Design . . . 27
3.6.2 Implementation . . . 28

4 Simulation Results **30**

5 Fractal Dimension of DLA **31**
5.1 Introduction to Dimension . . . 31
5.1.1 Hausdorff Dimension . . . 31
5.1.2 Box-Counting Dimension (D_{box}) . . . 32
5.1.3 Growth Rate Dimension (D_β) . . . 32
5.1.4 The Link Between D_{box} and D_β . . . 33
5.2 Methodology . . . 34
5.2.1 Calculating Box Dimension: The Box Size Problem . . . 34
5.2.2 Calculating Growth Rate Dimension: Sampling and Uncertainty . . . 35
5.3 Results and Statistical Errors . . . 36
5.4 Comparing the Evolution of Dimension and Growth Rate . . . 36

6 Anisotropy in DLA **39**
6.1 What is Anisotropy? . . . 39
6.2 Angular Radius of Gyration . . . 39
6.3 Angular Growth Rate Analysis . . . 39
6.4 Conclusions . . . 39

A Implementation Details and Optimisations **42**
A.1 Computational Framework . . . 42
A.1.1 The Numba JIT Compiler . . . 42
A.1.2 Rounding Semantics . . . 42
A.2 Spatial Indexing and Hierarchical Skipping . . . 43
A.2.1 Lattice: Flattened Bit-Map . . . 43
A.2.2 Off-Lattice: Dual-Structure Strategy . . . 44
A.2.3 2. The Spatial Hash (Hitting) . . . 44
A.3 Rejection Sampling . . . 44
A.3.1 Concept and Motivation . . . 44
A.3.2 The Algorithm . . . 45
A.3.3 Application 1: The Fuzzy Annulus . . . 45
A.3.4 Application 2: Walk-to-Line . . . 45
A.4 Walker's Alias Method . . . 46
A.4.1 Motivation . . . 46
A.4.2 Implementation . . . 46
A.5 Numerical Approximations . . . 47
A.5.1 Custom Bessel Functions . . . 47

A.6 Usability and Data Management 47
A.6.1 Heuristic Grid Sizing 47
A.6.2 Standardised Coordinate Output 47

Chapter 1

Introduction

1.1 Overview

Diffusion Limited Aggregation (DLA) is a stochastic model for irreversible colloidal aggregation. In this model a seed particle is placed at the origin, and then subsequent particles are released from a boundary far from the cluster one at a time. The particle then diffuses, either as a random walk on a lattice or Brownian motion (the continuous analogue to a random walk). Upon collision, the walker sticks and the next one starts diffusing.

The model was first proposed in 1981 by Witten and Sander to model the colloidal aggregation of metal precipitates [1], but has since been found to have applications in a diverse range of physical systems, including electrodeposition [2], bacteria growth [3] and viscous fingering in Hele-Shaw cells [4] among others.

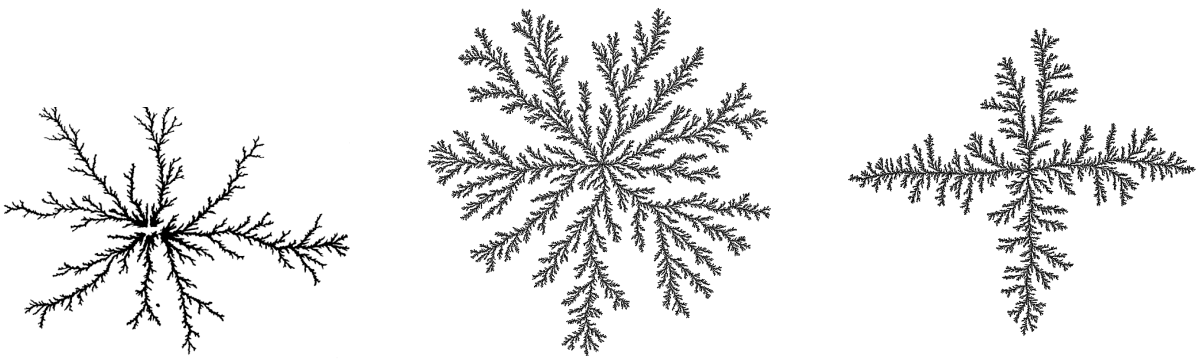

(a) Zinc Electrodeposition [2] **(b)** Off-Lattice DLA **(c)** On-Lattice DLA

Figure 1.1: Comparison of physical aggregation (a) with the two primary modelling approaches: continuum Off-Lattice DLA (b) and discrete On-Lattice DLA (c). Both are the results of simulation with 10^7 particles. Note the distinct grid-induced anisotropy visible in the on-lattice case.

1.2 Theoretical Context

While DLA is defined by a stochastic algorithm (random walkers), its theoretical significance lies in its connection to potential theory, the study of harmonic functions. As will be detailed in the following chapters, the probability density $\phi(u)$ of a diffusing particle at position u satisfies the Laplace equation $\nabla^2\phi = 0$. This is because both random walks and harmonic functions

are defined by the property that the value at any point is the average of its neighbours. Each additional walker to aggregate onto the cluster distorts this probability field, creating a Laplacian instability that amplifies irregularities into the observed complicated structure [5].

1.3 Fractal Dimension

Visually, DLA clusters exhibit self-similarity, where small branches look statistically identical to large branches. Mathematically this can be understood via fractal dimension D. This will be formally defined later, but can be thought of as the relationship between the mass of a cluster (number of particles N) and its linear size R, where:

$$N \propto R^D. \tag{1.1}$$

For instance, doubling the size of a line would double its mass, so it has dimension 1, while doubling the size of a solid disc would quadruple its mass, so this has dimension 2. Simulations of off-lattice DLA have consistently yielded a fractal dimension $D \approx 1.71$, but this has not been proven [6].

In fact, the only mathematical theorem on 2D DLA is due to Kesten, who proved a bound on the growth rate of on-lattice DLA in 1987, with a key consequence being that the fractal dimension of a DLA cluster is bounded below by 1.5 [7].

1.4 Anisotropy and Thesis Goals

It is likely that the dimension of on-lattice DLA is actually lower than off-lattice due to the effect of the grid on the cluster shape, a phenomenon known as anisotropy. Grebenkov and Beliaev discussed how anisotropy is likely to reduce the dimension of large enough clusters [8], but understanding why and to what extent is still not fully understood.

To answer these questions, simulations of DLA must be done both on and off-lattice. Additionally, to understand the effect of the lattice we create a 'hybrid' algorithm with continuous diffusion but a discrete aggregation method. Other hybrid models of DLA have previously been studied [9], including a similar semi-lattice approach [10], but to our knowledge have not been utilised to decouple the effects of transport from aggregation in on-lattice DLA.

The purpose of this dissertation is to compare these three different models of DLA, in order to decouple the diffusion process of DLA from the geometry of the aggregation. By analysing various features of the resulting clusters, we hope to better understand the relationship between the continuous and discrete variations of the probability distribution governing DLA growth.

We begin with a brief outline of the mathematical necessary to understand the underlying mechanisms behind DLA in chapter 2, before following on by explaining how to utilise this theory to create efficient algorithms, as well as discussing our hybrid approach, which forms chapter 3. In ?? we begin analysis, focusing on ???

Chapter 2

Potential Theory

2.1 Preliminaries

2.1.1 Martingales and the Optional Stopping Theorem

In this section, we review the necessary tools in probability for linking random walks to potential theory. We assume standard measure-theoretic knowledge (equivalent to the Part A Integration course). We begin by defining a probability space.

Definition 2.1 (Probability Space)**.** *A probability space* $(\Omega, \mathcal{F}, \mathbb{P})$ *is a triple where* Ω *is a sample set,* $\mathcal{F}$ *is a* σ*-algebra on* Ω *representing the set of events, and* $\mathbb{P}\colon \mathcal{F} \to [0,1]$ *is a countably additive set function such that* $\mathbb{P}(\Omega) = 1$*.*

We now introduce filtrations, which formalise the idea of accumulating information over time.

Definition 2.2 (Filtration)**.** *A filtration on a probability space* $(\Omega, \mathcal{F}, \mathbb{P})$ *is a sequence of sub-*σ*-algebras* $(\mathcal{F}_n)_{n\geq 0}$ *such that:*

$$\mathcal{F}_0 \subseteq \mathcal{F}_1 \subseteq \mathcal{F}_2 \subseteq \cdots \subseteq \mathcal{F}.$$

The tuple $(\Omega, \mathcal{F}, (\mathcal{F}_n)_{n\geq 0}, \mathbb{P})$ *is called a filtered probability space.*

We now consider how decisions can be made at a given time based on the information available up to that time.

Definition 2.3 (Stopping Time)**.** *A random variable* $\tau\colon \Omega \to \{0, 1, \ldots, \infty\}$ *is called a stopping time with respect to the filtration* $(\mathcal{F}_n)_{n\geq 0}$ *if for every* $n \geq 0$*:*

$$\{\omega \in \Omega : \tau(\omega) \leq n\} \in \mathcal{F}_n.$$

For example, “stop when the walker hits the cluster" is a valid stopping time but “stop at the step before the walker hits the cluster" is not.
With this in place, we model an unbiased process via the concept of a martingale.

Definition 2.4 (Martingale)**.** *A discrete-time stochastic process* $(X_n)_{n\geq 0}$ *is a martingale with respect to a filtration* $(\mathcal{F}_n)_{n\geq 0}$ *if:*

1. X_n *is adapted to the filtration (i.e.,* X_n *is* $\mathcal{F}_n$*-measurable for all* n*),*

2. $\mathbb{E}[|X_n|] < \infty$ *for all* n,

3. $\mathbb{E}[X_{n+1} \mid \mathcal{F}_n] = X_n$ *almost surely.*

A martingale essentially encapsulates the idea of a "fair game". The condition $\mathbb{E}[X_{n+1} \mid \mathcal{F}_n] = X_n$ implies that the best prediction for the next value is the current value, even with the entire history of the process.
A key result regarding martingales is the Optional Stopping Theorem. There are multiple versions, and we only state one here sufficient for our purposes.

Theorem 2.1 (Doob's Optional Stopping Theorem)**.** *Let X be a martingale on a filtered probability space. If τ is a stopping time such that $\mathbb{P}(\tau < \infty) = 1$, and X is uniformly bounded (i.e., there exists K such that $|X_n(\omega)| \leq K$ for all n, ω), then:*

$$\mathbb{E}[X_\tau] = \mathbb{E}[X_0]. \tag{2.1}$$

Proof. See Williams [11, Section 10.10]. □

This means that no stopping strategy τ can change the expected payout of a fair game. This will allow us to relate the value of a harmonic function at a starting point x (start of the walk) to its expected value at the boundary ∂D (end of the walk).

2.1.2 Measure Theory Concepts

In the study of potential theory, we will want to understand the relationship between the geometric length of a boundary (Lebesgue measure) and the probability that a random walker hits that boundary (harmonic measure). To rigorously define a probability density function linking these two concepts we require the notion of the Radon-Nikodym derivative.

We begin by defining absolute continuity between two measures, which is a necessary prerequisite. *Remark:* throughout this work, we assume our measurable spaces are equipped with the Borel σ-algebra.

Definition 2.5 (Absolute Continuity)**.** *Let $(\Omega, \mathcal{F})$ be a measurable space equipped with two measures μ and ν. We say that ν is absolutely continuous with respect to μ, denoted $\nu \ll \mu$, if for every $A \in \mathcal{F}$:*

$$\mu(A) = 0 \implies \nu(A) = 0. \tag{2.2}$$

Theorem 2.2 (The Radon-Nikodym Theorem)**.** *Let $(\Omega, \mathcal{F})$ be a measurable space, and let μ and ν be two finite measures such that $\nu \ll \mu$. Then, there exists a non-negative measurable function $f \colon \Omega \to [0, \infty)$, unique μ-almost everywhere, such that for any $A \in \mathcal{F}$:*

$$\nu(A) = \int_A f \, d\mu. \tag{2.3}$$

The function f is called the Radon-Nikodym derivative of ν with respect to μ and is denoted by:

$$f = \frac{d\nu}{d\mu}. \tag{2.4}$$

Proof. See Williams [11, Section 14.13]. □

2.1.3 Conformal Maps

Here we state the definition of a conformal map, and the Riemann mapping theorem, both from Ahlfors [12]. Conformal maps will be very important for two reasons: theoretically, the Riemann Mapping Theorem guarantees that the Dirichlet problem we consider later in this chapter is well-posed even on the potentially fractal boundary of a DLA cluster. Algorithmically, these maps allow us to exploit the 'conformal invariance' of Brownian motion, which we discuss in the latter parts of this chapter and into the next.

Definition 2.6 (Conformal Map)**.** *Let $U, V \subseteq \mathbb{C}$ be open sets. A function $f\colon U \to V$ is called* ***conformal*** *at a point $z_0 \in U$ if it preserves angles and orientation locally.*

Analytically, a holomorphic function f is conformal at every point where its derivative is non-zero, i.e., $f'(z) \neq 0$. If f is a holomorphic bijection between U and V (a biholomorphism), it is essentially a global conformal map, and we say U and V are conformally equivalent.

Theorem 2.3 (Riemann Mapping Theorem)**.** *Let $U \subsetneq \mathbb{C}$ be a non-empty, simply connected open subset of the complex plane. Then, there exists a unique bijective conformal map f from U onto the open unit disk $\mathbb{D} = \{z \in \mathbb{C} : |z| < 1\}$ such that:*

$$f(z_0) = 0 \quad \textit{and} \quad f'(z_0) > 0, \tag{2.5}$$

for any prescribed point $z_0 \in U$.

Proof. See Ahlfors [12, Chapter 6.1]. □

2.2 Discrete Potential Theory[1]

We now introduce potential theory in a discrete setting, focusing on its applications to DLA. We start by considering the definitions of a random walk and a discrete harmonic function.

2.2.1 Simple Random Walks and the Generator

Definition 2.7 (Random Walk)**.** *A random walk on $\mathbb{Z}^2$ is a stochastic process $(S_n)_{n\geq 0}$ defined by $S_n = S_0 + \sum_{i=1}^{n} X_i$, where $X_1, X_2, \ldots$ are independent and identically distributed random variables taking values in $\mathbb{Z}^2$.*

The common distribution of the steps is denoted by $p(x) = \mathbb{P}(X_i = x)$ and is called the ***increment distribution****.*

For the purposes of DLA, we are primarily interested in the symmetric case:

Definition 2.8 (Simple Symmetric Random Walk)**.** *A random walk is called a Simple Symmetric Random Walk (SSRW) if the increment distribution is uniform on the nearest neighbours:*

$$p(x) = \begin{cases} \frac{1}{4} & \textit{if } |x| = 1, \\ 0 & \textit{otherwise.} \end{cases} \tag{2.6}$$

[1] The theoretical framework in this section is largely adapted from Lawler and Limic [5], applied here to the specific case of 2D DLA.

We now define the operator that governs the evolution of the expectation of functions under this flow.

Definition 2.9 (Discrete Generator)**.** *Let $p(\cdot)$ be the increment distribution of a random walk. The generator $\mathcal{L}$ is a linear operator acting on functions $f : \mathbb{Z}^2 \to \mathbb{R}$ defined by:*

$$\mathcal{L}f(x) = \sum_{y \in \mathbb{Z}^2} p(y - x)[f(y) - f(x)]. \tag{2.7}$$

Probabilistically, this represents the expected change in the function f in a single step of the walk starting at x:

$$\mathcal{L}f(x) = \mathbb{E}[f(S_1) - f(S_0) \mid S_0 = x]. \tag{2.8}$$

Remark: For SSRW, the generator is proportional to the (unnormalised) discrete Laplacian, $\mathcal{L} = \frac{1}{4}\Delta$, where $\Delta f(x) = \sum_{y \sim x}(f(y) - f(x))$.

Definition 2.10 (Discrete Harmonic Functions)**.** *A function $f : A \cup \partial A \to \mathbb{R}$ is said to be harmonic in a subset $A \subset \mathbb{Z}^2$ if:*

$$\mathcal{L}f(x) = 0 \quad \textit{for all } x \in A. \tag{2.9}$$

Proposition 2.1 (Harmonic Functions are Martingales)**.** *Let S_n be a random walk and f be a function harmonic on a set A. Let $\tau_A = \inf\{n \geq 0 : S_n \notin A\}$ be the first exit time from A. Then the stopped process $M_n = f(S_{n \wedge \tau_A})$ is a martingale (with respect to the natural filtration of S_n).*

Proof. See Lawler and Limic [5, Proposition 6.1.1]. □

2.2.2 The Discrete Dirichlet Problem

Definition 2.11 (Boundary)**.** *Before proceeding, we distinguish between the occupied cluster and the available growth sites. Let $K \subset \mathbb{Z}^2$ denote the set of occupied sites (the cluster). We define the boundary ∂K as the outer boundary consisting of all empty sites adjacent to the cluster:*

$$\partial K = \{y \in \mathbb{Z}^2 \setminus K : \exists x \in K \textit{ such that } |y - x| = 1\}. \tag{2.10}$$

Subsequently we refer to the outer boundary as the ***boundary*** *of the cluster ∂K.*

We now attempt to model DLA as a discrete Dirichlet problem, where the unknown function $u_R(x)$ is defined as the probability of a random walker starting at x escaping to some outer boundary ∂R of radius R before hitting the cluster boundary ∂K.

$$u_R(x) = \mathbb{P}^x(\text{hit outer boundary } \partial R \text{ before } \partial K). \tag{2.11}$$

As the diffusion process is a simple random walk, the probability of escaping from x is simply the average of the excape probabilities of its neighbours:

$$u_R(x) = \frac{1}{4} \sum_{z \sim x} u_R(z). \tag{2.12}$$

By definition this makes $u(\cdot)$ harmonic, hence satisfies the discrete Laplace equation $\mathcal{L}u(x) = 0$.

Thus, each stage of DLA is equivalent to solving the Dirichlet problem:

$$\begin{cases} \Delta f(x) = 0 & \text{if } x \notin K \text{ (Harmonic in free space)}, \\ f(x) = 0 & \text{if } x \in \partial K \text{ (Hit the cluster)}, \\ f(x) = 1 & \text{if } x \in \partial R \text{ (Escaped to infinity)}. \end{cases} \tag{2.13}$$

Remark 2.1. *We could have solved for the probability of hitting the cluster before escaping, reversing the boundary conditions. However, we do it this way to preserve consistency with the later sections and the electrostatic interpretation (see Appendix).*

We now look at the general discrete Dirichlet problem in 2D:

Theorem 2.4 (Solution to the 2D Discrete Dirichlet Problem)**.** *Let $A \subset \mathbb{Z}^2$ be a subset (potentially infinite) and let $\tau = \inf\{n \geq 0 : S_n \notin A\}$ be the first exit time. Assume that:*

1. *The random walk stops almost surely: $\mathbb{P}_x(\tau < \infty) = 1$ for all $x \in A$.*
2. *The boundary function $g : \partial A \to \mathbb{R}$ is bounded.*

Then, the function $f : A \cup \partial A \to \mathbb{R}$ defined by the expectation:

$$f(x) = \mathbb{E}_x[g(S_\tau)] \tag{2.14}$$

is the unique bounded solution to the Dirichlet system:

$$\mathcal{L}f(x) = 0, \qquad \text{for all } x \in A, \tag{2.15a}$$

$$f(x) = g(x), \qquad \text{for all } x \in \partial A. \tag{2.15b}$$

Proof. **Existence:** We first check that (2.14) satisfies the system (2.15). If $x \in \partial A$, then $\tau = 0$, so $S_\tau = x$ and $f(x) = \mathbb{E}_x[g(x)] = g(x)$. Thus the boundary condition (2.15b) holds. If $x \in A$, we use the Markov property at the first step:

$$f(x) = \mathbb{E}_x[g(S_\tau)] = \sum_{y \sim x} p(y - x)\mathbb{E}_y[g(S_\tau)] = \sum_{y \sim x} p(y - x)f(y). \tag{2.16}$$

Thus $\mathcal{L}f(x) = 0$, satisfying (2.15a).

Uniqueness: Suppose h is another bounded solution to the system (2.15). Since h satisfies (2.15a), it is harmonic, and by proposition 2.1, the process $M_n = h(S_{n\wedge\tau})$ constitutes a martingale. Since h is bounded, M_n is a bounded martingale.
Given that $\tau < \infty$ almost surely, we can apply the Optional Stopping Theorem (theorem 2.1):

$$h(x) = \mathbb{E}_x[M_0] = \mathbb{E}_x[M_\tau] = \mathbb{E}_x[h(S_\tau)]. \tag{2.17}$$

From condition (2.15b), we know $h(S_\tau) = g(S_\tau)$ on the boundary. Therefore:

$$h(x) = \mathbb{E}_x[g(S_\tau)], \tag{2.18}$$

which is exactly $f(x)$. Thus $h \equiv f$, proving uniqueness. □

2.2.3 Recurrence and the Potential Kernel

To apply this solution to DLA, we require taking the limit $R \to \infty$, which is equivalent to solving the Dirichlet problem with a boundary condition at infinity.

This is done formally by Lawler and Limic [5, Proposition 6.2.2], who prove that a solution to the Dirichlet problem in d dimensions is uniquely determined by values on the finite boundary ∂A and the value at infinity, taking the form:

$$f(x) = \mathbb{E}_x[g(S_\tau)\mathbf{1}_{\tau<\infty}] + g(\infty)\mathbb{P}_x(\tau = \infty). \tag{2.19}$$

However, due to the recurrence of a random walk in 2D space, the escape probability $\mathbb{P}_x(\tau = \infty)$ is zero for any random walk in $\mathbb{Z}^2$. Consequently, the second term vanishes. The solution is now entirely determined by the inner boundary ∂K. For our boundary conditions given in eq. (2.15), the only solution is the trivial potential $f \equiv 0$.

To address this we now turn to the concept of Green's functions. The Green's function $G(x, y)$ represents the potential field at x induced by a point source at y. Put another way, it is the solution to the Poisson equation:

$$\Delta G(x, y) = -\delta_y(x). \tag{2.20}$$

In the probabilistic setting, the Green's function has a precise interpretation: it is the expected number of times a random walk starting at x visits the site y:

$$G(x, y) = \mathbb{E}_x\left[\sum_{n=0}^{\infty} \mathbf{1}_{\{S_n=y\}}\right] = \sum_{n=0}^{\infty} p_n(y - x). \tag{2.21}$$

For transient random walks (such as in dimensions $d \geq 3$), this sum converges, and $G(x, y)$ decays to zero as $|x - y| \to \infty$.

However, in two dimensions the random walk is recurrent, meaning a walker returns to the origin infinitely often, causing the sum in the definition of $G(x, y)$ to diverge:

$$G(x, y) = \infty \quad \text{for all } x, y. \tag{2.22}$$

Instead we regularise the divergent Green's function by working with potential differences, motivating the definition of the potential kernel.

Definition 2.12 (Potential Kernel)**.** *The potential kernel $a(x)$ for simple random walk in $\mathbb{Z}^2$ is defined as:*

$$a(x) = \sum_{n=0}^{\infty}[p_n(0) - p_n(x)], \tag{2.23}$$

where $p_n(x) = \mathbb{P}(S_n = x)$.

It is not immediately clear that this will be finite. However, using estimates from the Local Central Limit Theorem [5, Section 4.4.1], it can be shown that for any fixed x, the difference decays as:

$$|p_n(0) - p_n(x)| \leq C|x|n^{-3/2}. \tag{2.24}$$

Since the series $\sum_{n=1}^{\infty} n^{-3/2}$ converges absolutely, the potential kernel $a(x)$ is well-defined and

finite for all x.

By transitioning from the Dirichlet problem to the potential kernel, we stop solving directly for the probability and instead solve for a renormalised potential $a(x)$. Intuitively, $a(x)$ quantifies the relative accessibility of a given lattice site from infinity. Importantly, the growth probability is recovered through the discrete differences of this potential: The proportion of walkers hitting a boundary point depends on the 'steepness' of the potential at the boundary.
The potential kernel therefore acts as the fundamental solution to the Laplacian on the lattice. It satisfies the Poisson equation:

$$\mathcal{L}a(x) = \delta_0(x) = \begin{cases} 1 & x = 0, \\ 0 & x \neq 0. \end{cases} \tag{2.25}$$

It replaces the 'boundary at infinity' with an asymptotic growth condition.

Theorem 2.5 (Asymptotics of the Potential Kernel)**.** *As $|x| \to \infty$, the potential kernel satisfies:*

$$a(x) = \frac{2}{\pi} \ln |x| + k + O(|x|^{-2}), \tag{2.26}$$

where k is a constant involving Euler's constant γ.

Proof. See Lawler and Limic [5, Theorem 4.4.4]. □

Remark 2.2. *The logarithmic growth may seem arbitrary, but can in fact be justified with some geometric intuition: Consider a circle of radius R centred on the cluster (large enough so that it contains the entire cluster). For a steady flow of random walkers coming from infinity to reach the cluster, the total flux (rate of flow) across this circle must be independent of R. Else walkers would have to spontaneously appear or disappear which does not fit the model. The total flux is the product of the circumference ($2\pi R$) and the radial gradient of the potential ($\nabla f(R)$). This implies that the gradient must decay as $\nabla f(R) \propto 1/R$. Integrating this gradient yields:*

$$f(R) \propto \int \frac{1}{R}\, dR = \ln R. \tag{2.27}$$

2.2.4 Discrete Harmonic Measure

We now formally define the discrete harmonic measure:

Definition 2.13 (Harmonic Measure from Infinity)**.** *Let $K \subset \mathbb{Z}^2$ be a finite cluster. The harmonic measure $H_{\partial K}(y)$ is the probability that a simple random walk starting from infinity first hits the cluster K at the point $y \in \partial K$. Formally, it is defined as the limit:*

$$H_{\partial K}(y) = \lim_{|x| \to \infty} \mathbb{P}_x(S_{\tau_K} = y), \tag{2.28}$$

where $\tau_K = \inf\{n \geq 0 : S_n \in K\}$ is the first hitting time.

Lawler and Limic [5, Proposition 6.6.1] prove that this limit exists and defines a unique probability measure on the boundary ∂K. While the definition may seem abstract, we can in fact calculate these probabilities directly from the potential field derived in the previous section.

We begin by formally defining the concept of discrete flux which was invoked in remark 2.2. On the lattice, the concept of a normal derivative $\nabla u \cdot \mathbf{n}$ is replaced by a difference operator across the boundary.

Definition 2.14 (Discrete Flux)**.** *For a function $u : \mathbb{Z}^2 \to \mathbb{R}$ that vanishes on K, the discrete flux at a boundary point $y \in \partial K$ is the sum of the potential differences from its neighbours in the complement:*

$$\Phi(y) = \sum_{z \sim y, z \notin K} (u(z) - u(y)). \tag{2.29}$$

Since $u(y) = 0$ on the boundary, this simplifies to the sum of the external neighbours: $\Phi(y) = \sum_{z \sim y, z \notin K} u(z)$.

A fundamental result that we require from discrete potential theory is the equivalence of flux and the harmonic measure.

Proposition 2.2 (Harmonic Measure as Flux)**.** *The harmonic measure at $y \in \partial K$ is proportional to the discrete flux at y:*

$$H_{\partial K}(y) = \frac{\Phi(y)}{\sum_{w \in \partial K} \Phi(w)}. \tag{2.30}$$

Proof. For a formal rigorous proof connecting hitting probabilities to the Laplacian of the Green's function, see Lawler and Limic [5, Proposition 1.6.3]. □

To understand this link conceptually, we interpret the potential $u(x)$ as the density of a stream of random walkers originating from infinity and disappearing upon hitting the cluster K. In the exterior free space, the condition $\mathcal{L}u(x) = 0$ implies a sort of conservation law: the expected number of walkers entering any site x equals the number leaving it. However, at a boundary site $y \in \partial K$, walkers arrive from neighbours but do not continue; they form the cluster. The operator $\mathcal{L}u(y)$ represents this absorption rate:

$$\text{Rate into } y \propto \sum_{z \sim y} (\text{Density at } z) = \sum_{z \sim y} u(z) = \Phi(y). \tag{2.31}$$

Since the harmonic measure $H(y)$ is the proportion of all walkers that terminate at y, it is directly proportional to this local absorption rate, i.e., the flux $\Phi(y)$.

This concludes the discrete formulation of the DLA process. The growth mechanism can be summarised as follows:

1. The system is governed by a harmonic potential $u(x)$, which is determined by the geometry of the cluster K.

2. This potential is the solution to the discrete Laplace equation $\mathcal{L}u = 0$ with $u|_K = 0$ and $u(x) \sim \ln |x|$ as $|x| \to \infty$.

3. The growth probability at a boundary site y is determined by the local flux $\Phi(y) = \sum_{z \sim y} u(z)$, where a new particle is added with probability $H(y) \propto \Phi(y)$.

2.3 Continuous Potential Theory

[2]

We now provide an introduction to potential theory in the continuous setting. Much of this section is formally derived by extending section 2.2 to the continuous case using stochastic calculus. However, for our purposes, an understanding of the main results is sufficient to build a robust intuition for DLA, so we focus only on the useful results. The fundamental change is replacing the discrete random walker with its scaling limit: Brownian motion.

2.3.1 Brownian Motion and the Laplacian

We start by defining Brownian motion in both $\mathbb{R}^d$ and $\mathbb{C}$.

Definition 2.15 (Standard Brownian Motion)**.** *A standard Brownian motion in $\mathbb{R}^d$ is a stochastic process $B = (B_t)_{t\geq 0}$ defined on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ taking values in $\mathbb{R}^d$, such that:*

1. *$B_0 = 0$ almost surely.*
2. *The process has independent increments: for any $0 \leq s < t$, $B_t - B_s$ is independent of $\mathcal{F}_s = \sigma(B_u : u \leq s)$.*
3. *The increments are Gaussian: $B_t - B_s \sim \mathcal{N}(0, (t-s)I_d)$.*
4. *The paths $t \mapsto B_t$ are continuous almost surely.*

In the study of 2D DLA, we will utilise results from complex analysis. Thus, it is natural to identify the plane $\mathbb{R}^2$ with $\mathbb{C}$.

Definition 2.16 (Complex Brownian Motion)**.** *A canonical complex Brownian motion Z_t is a process taking values in $\mathbb{C}$, defined as:*

$$Z_t = X_t + iY_t, \tag{2.32}$$

where X_t and Y_t are two independent, standard one-dimensional Brownian motions.

The utility of this complex formulation will be apparent in section 2.3.3.

Remark 2.3. *We gloss over much of the detail in this section, as for our purposes a complete rigourous introduction to stochastic calculus is unnecessary. Intuitively, Brownian motion can be understood as the continuous limit of a random walk. In fact, this has been rigorously proven as Donsker's Invariance Principle (see Mörters and Peres [13, Theorem 5.22]). For DLA, this implies that large clusters formed by random walks should exhibit the same macroscopic properties as those formed by Brownian particles. For a full treatment of stochastic calculus and Brownian motion, see Course B8.2 or Mörters and Peres [13].*

The connection between the physical process (Brownian motion) and the geometric potential (Laplace's equation) is provided by the infinitesimal generator.

[2] Again, this is all existing theory, we merely synthesise results regarding Brownian motion and potential theory given in Mörters and Peres [13] for the context of DLA.

Definition 2.17 (Continuous Generator and Laplacian)**.** *The Laplace operator (or Laplacian) Δ in $\mathbb{R}^d$ is the sum of the second partial derivatives:*

$$\Delta = \sum_{i=1}^{d} \frac{\partial^2}{\partial x_i^2}. \tag{2.33}$$

The infinitesimal generator of standard Brownian motion is the operator $\mathcal{A} = \frac{1}{2}\Delta$. It characterises the process via the martingale property: for any compactly supported C^2 (twice continuously differentiable) function f, the quantity

$$M_t = f(B_t) - f(B_0) - \int_0^t \frac{1}{2}\Delta f(B_s)\, ds \tag{2.34}$$

is a martingale with respect to the filtration of B_t.

The integral term in definition 2.17 represents the expected drift of the function along the path. By subtracting this drift, we isolate the centred random noise, which is a martingale.

Remark 2.4 (Scaling Factors and Independence)**.** *We point out here that the scaling factor here differs from that in section 2.2.1. This discrepancy stems from the second-order term of the Taylor expansion used to derive the Brownian limit. However, since the potential distribution for DLA is determined by the condition $\mathcal{L}u = 0$ (or $\Delta u = 0$), this constant factor cancels out. The resulting harmonic measure is invariant. For a full derivation of this limit, see Lawler and Limic [5, Section 2.1].*

Definition 2.18 (Harmonic Functions)**.** *A twice continuously differentiable function $u : D \to \mathbb{R}$ defined on a domain $D \subset \mathbb{R}^d$ is called harmonic if it lies in the kernel of the Laplace operator:*

$$\Delta u(x) = 0 \quad \text{for all } x \in D. \tag{2.35}$$

By comparing this with the generator definition above, we see immediately that if u is harmonic, the drift term vanishes. This yields the critical result that $u(B_t)$ is a martingale, implying that the "average value" of a harmonic function is conserved along Brownian paths. We formalise this as the mean value property.

Theorem 2.6 (Mean Value Property)**.** *A locally integrable function $u : D \to \mathbb{R}$ is harmonic ($\Delta u = 0$) if and only if for every ball $B(x, r)$ strictly contained in the domain D:*

$$u(x) = \frac{1}{|\partial B(x,r)|} \int_{\partial B(x,r)} u(y)\, d\sigma(y), \tag{2.36}$$

where σ denotes the surface measure on the boundary sphere.

This is the continuous equivalent to the 'average of its neighbours' property of in the discrete setting.

2.3.2 Harmonic Measure

Before we present the continuous analogue to theorem 2.4, we define the continuous version of the harmonic measure.

Definition 2.19 (Harmonic Measure)**.** *Let $D \subset \mathbb{R}^d$ be a domain. The harmonic measure $\omega_x = \omega_x(D, \cdot)$ is the probability measure on the boundary ∂D defined by:*

$$\omega_x(E) = \mathbb{P}^x(B_{\tau_D} \in E), \tag{2.37}$$

for Borel sets $E \subset \partial D$, where $\tau_D = \inf\{t > 0 : B_t \notin D\}$ is the first exit time.

The solution to the continuous Dirichlet problem is given by Kakutani's Theorem.

Theorem 2.7 (Kakutani's Theorem)**.** *If D is a bounded domain and g is a continuous function on ∂D, then the function defined by:*

$$u(x) = \mathbb{E}^x[g(B_{\tau_D})] = \int_{\partial D} g(y)\, d\omega_x(y) \tag{2.38}$$

is harmonic in D and satisfies $\lim_{x \to y} u(x) = g(y)$ for regular boundary points $y \in \partial D$. Here, ω_x denotes the harmonic measure relative to x, representing the distribution of the first hitting point of Brownian motion on ∂D.

Proof. For a rigorous proof relying on the martingale property of harmonic functions, see Mörters and Peres [13, Theorem 3.12]. □

As in the discrete case, we are concerned with extending this to infinity. As a consequence of the following lemma, 2D Brownian motion is neighbourhood recurrent, meaning it will revisit any neighbourhood infinitely often. This prevents us from simply imposing a boundary condition at infinity, just as recurrence did in the discrete case.

Lemma 2.1 (Exit Probabilities from an Annulus)**.** *Let B_t be a complex Brownian motion starting at z_0 within the annulus $A = \{z \in \mathbb{C} : r < |z| < R\}$. Let $T_A = \inf\{t > 0 : B_t \notin A\}$ be the first exit time. The probability that the process exits through the inner boundary $|z| = r$ is given by:*

$$\mathbb{P}^{z_0}(|B_{T_A}| = r) = \frac{\ln R - \ln |z_0|}{\ln R - \ln r}. \tag{2.39}$$

Proof. See Mörters and Peres [13, Theorem 3.18]. The proof uses Itô's formula (stochastic chain rule) on the function $\ln |Z_t|$. □

Corollary 2.1 (Recurrence)**.** *Planar Brownian motion is neighbourhood recurrent. For any starting point z_0 and any $\epsilon > 0$, the process will enter the disk $\{z : |z - z_0| < \epsilon\}$ almost surely:*

$$\mathbb{P}(\exists t > 0 : |B_t - z_0| < \epsilon) = 1. \tag{2.40}$$

Proof. Taking the limit $R \to \infty$ in the annulus exit probability yields $\mathbb{P}(|B_{T_A}| = r) \to 1$. □

2.3.3 Conformal Invariance

While in the discrete case we resolved this recurrence issue by turning to (regularised) Green's functions, we will instead utilise a fundamental property of Brownian motion: conformal invariance.

Theorem 2.8 (Lévy's Conformal Invariance Theorem)**.** *Let $U \subset \mathbb{C}$ be a domain and $f : U \to V$ be a non-constant analytic function. If Z_t is a complex Brownian motion in U, then the image process $W_t = f(Z_t)$ is, up to a random time change, a complex Brownian motion in V. Specifically, if $\tau_t = \int_0^t |f'(Z_s)|^2 ds$, then $f(Z_{\tau^{-1}(t)})$ is a standard Brownian motion in V.*

Proof. See Mörters and Peres [13, Theorem 7.20]. □

Corollary 2.2 (Conformal Invariance of Harmonic Measure)**.** *Let $f : D \to D'$ be a conformal map. The harmonic measure is invariant under f:*

$$\omega_z(D, E) = \omega_{f(z)}(D', f(E)) \tag{2.41}$$

for any Borel set $E \subset \partial D$.

This result guarantees that utilising conformal maps in the sampling process will not distort the underlying harmonic measure, preserving the integrity of our off-lattice algorithm in the subsequent section.

2.3.4 Solvable Geometries

We conclude by introducing the probability density function for the hitting probabilities, and finding it explicitly for two fundamental domains used in DLA algorithms: the Upper Half-Plane ($\mathbb{H}$) and the Unit Disk ($\mathbb{D}$).

Definition 2.20 (Poisson Kernel)**.** *Let σ be the surface measure on ∂D. If the harmonic measure is absolutely continuous with respect to σ, the Poisson kernel $P(x, y)$ is defined as the Radon-Nikodym derivative:*

$$P(x, y) = \frac{d\omega_x}{d\sigma}(y). \tag{2.42}$$

It represents the probability density of the Brownian motion exiting D at y given it started at x.

We can now express the conformal invariance of the harmonic measure in terms of the Poisson kernel. By definition, the harmonic measure of a set $E \subset \partial D$ is the integral of the Poisson kernel with respect to the surface measure $d\sigma$:

$$\omega_z(D, E) = \int_E P_D(z, \zeta) d\sigma(\zeta). \tag{2.43}$$

Applying the conformal invariance corollary to a map $f : D \to \mathbb{D}$ such that $f(z) = 0$, and noting that the Poisson kernel on the disk from the origin is uniform ($P_{\mathbb{D}}(0, w) = \frac{1}{2\pi}$), we have:

$$\int_E P_D(z, \zeta) d\sigma(\zeta) = \omega_0(\mathbb{D}, f(E)) = \int_{f(E)} \frac{1}{2\pi} d\sigma(w). \tag{2.44}$$

By the change of variables $w = f(\zeta)$, the arc length element transforms as $d\sigma(w) = |f'(\zeta)| d\sigma(\zeta)$, resulting in:

$$\int_E P_D(z, \zeta) d\sigma(\zeta) = \int_E \frac{1}{2\pi} |f'(\zeta)| d\sigma(\zeta). \tag{2.45}$$

This equality holds for any Borel set E, hence:

$$P_D(z, \zeta) = \frac{1}{2\pi}|f'(\zeta)|. \tag{2.46}$$

We now apply this relation to our domains of interest:

Theorem 2.9 (Poisson Kernel for $\mathbb{H}$)**.** *Consider the upper half-plane $\mathbb{H} = \{z \in \mathbb{C} : Im(z) > 0\}$. The probability density of a Brownian particle starting at $z = x + iy$ hitting the real line at $u \in \mathbb{R}$ is given by the Cauchy distribution:*

$$P_{\mathbb{H}}(z, u) = \frac{1}{\pi}\frac{y}{(u-x)^2 + y^2}. \tag{2.47}$$

Proof. We construct a conformal map $\phi : \mathbb{H} \to \mathbb{D}$ that sends the starting point z to the origin 0, and the real line $\mathbb{R}$ to the unit disk. The standard choice is the Cayley transform variation:

$$\phi(w) = \frac{w - z}{w - \bar{z}}. \tag{2.48}$$

Observe that $\phi(z) = 0$, and for any real boundary point $u \in \mathbb{R}$, $|\phi(u)| = 1$. Using the pullback proposition, the kernel is given by $\frac{1}{2\pi}|\phi'(w)|$. We compute the derivative with respect to the w:

$$\phi'(w) = \frac{z - \bar{z}}{(w - \bar{z})^2}. \tag{2.49}$$

Now evaluated at a boundary point $w = u$:

$$|\phi'(u)| = \frac{|2iy|}{|u - \bar{z}|^2} = \frac{2y}{(u-x)^2 + y^2}. \tag{2.50}$$

Thus, the Poisson kernel is:

$$P_{\mathbb{H}}(z, u) = \frac{1}{2\pi}\left(\frac{2y}{(u-x)^2 + y^2}\right) = \frac{1}{\pi}\frac{y}{(u-x)^2 + y^2}. \tag{2.51}$$

□

We now follow the same procedure for the unit disk.

Theorem 2.10 (Poisson Kernel for $\mathbb{D}$)**.** *The Poisson kernel for the unit disk $\mathbb{D}$, representing the exit density at angle θ for a particle starting at $z \in \mathbb{D}$, is given by:*

$$P_{\mathbb{D}}(z, \theta) = \frac{1}{2\pi}\frac{1 - |z|^2}{|e^{i\theta} - z|^2}. \tag{2.52}$$

Proof. We seek an automorphism of the disk $\psi : \mathbb{D} \to \mathbb{D}$ that maps the starting point z to the origin 0. This is given by the Möbius transformation:

$$\psi(w) = \frac{w - z}{1 - \bar{z}w}. \tag{2.53}$$

The derivative at the boundary point w is:

$$\psi'(w) = \frac{1 - |z|^2}{(1 - \bar{z}w)^2}. \tag{2.54}$$

To find the magnitude $|\psi'(w)|$, we observe that for $w = e^{i\theta}$, the denominator satisfies:

$$|1 - \bar{z}e^{i\theta}| = |e^{-i\theta} - \bar{z}| = |\overline{e^{-i\theta} - \bar{z}}| = |e^{i\theta} - z|. \tag{2.55}$$

Therefore, the Poisson kernel is:

$$\frac{1}{2\pi}|\psi'(e^{i\theta})| = \frac{1}{2\pi}\frac{1 - |z|^2}{|e^{i\theta} - z|^2}. \tag{2.56}$$

□

Corollary 2.3 (Poisson Kernel for the Exterior of $\mathbb{D}$)**.** *Let $D_{ext} = \{z \in \mathbb{C} : |z| > 1\}$. The harmonic measure of the unit circle relative to a point $z \in D_{ext}$ is given by the density:*

$$P_{ext}(z, \theta) = \frac{1}{2\pi}\frac{|z|^2 - 1}{|z - e^{i\theta}|^2}. \tag{2.57}$$

Proof. Consider the geometric inversion $a = 1/\bar{z}$, which maps $D_{ext} \to \mathbb{D}$. By the conformal invariance of harmonic measure, the exit distribution for z is identical to that of a. Substituting a into the interior kernel (theorem 2.10) and simplifying yields the result. □

This concludes the continuous formulation of the DLA process. The theoretical framework parallels the discrete setting but utilises slightly different machinery:

1. The system is governed by a continuous potential field $u(x)$ determined by the geometry of the cluster, representing the probability density of the underlying Brownian motion.

2. The field $u(x)$ satisfies the continuous Laplace equation $\Delta u = 0$, but solving this explicitly is difficult. Instead we utilise the property of conformal invariance, allowing us to encode the potential field geometrically.

3. The growth probability at a boundary point is given by the Poisson Kernel, which serves as the continuous analogue to the discrete flux.

Chapter 3

Algorithms

3.1 Introduction

Simulating DLA presents a significant computational challenge. Directly simulating the random walk of each particle from large distance is prohibitively expensive. As the radius R of the cluster grows, the number of steps scales as $O(R^2)$. To generate large enough clusters for useful analysis, we employ accelerated algorithms that exploit the theory developed in chapter 2

In this chapter, we develop three distinct models:

1. **On-Lattice DLA:** A discrete implementation adapted from the algorithm by Loh [14].
2. **Off-Lattice DLA:** A continuous implementation using conformal mappings adapted from Bell [15].
3. **Hybrid DLA:** A novel model which combines continuous diffusion with on-lattice aggregation.

These models all share roughly the same structure, which we split into four stages: Initialisation, Far-Field Diffusion, Near-Field Diffusion and Aggregation. We will detail each of the stages for the first two models in the following sections. We will then present the hybrid model in section 3.6, which we implement as a variation of the off-lattice model.

3.2 Initialisation

The mathematical definition of DLA requires particles to originate "from infinity". In practice, we must approximate this by launching particles from a boundary $\partial\Omega$ sufficiently far from the cluster K. We now detail how this is achieved in both the off-lattice and on-lattice settings.

3.2.1 Off-Lattice: The Launching Circle

In the continuous setting, initialising at infinity is equivalent to initialising uniformly on a large bounding circle. To see this, let R_{max} be the radius of the smallest circle centred at the origin containing the cluster. We define a launching circle of radius $R_l > R_{max}$.

Now consider the conformal map $\phi(z) = \frac{R_l}{z}$. This creates a bijection between the exterior of the launching circle and its interior, mapping infinity to the origin. By theorem 2.6, the harmonic

measure of the origin in the disk is uniform (theorem 2.10 with $z = 0$). Consequently, the hitting distribution of a Brownian motion starting at infinity on the circle $|z| = R_l$ is uniform. Thus, our off-lattice algorithm simply samples a starting angle $\theta \sim U[0, 2\pi)$ and initialises the walker at $z_0 = R_l e^{i\theta}$. This is therefore exactly equivalent analytically to starting "at infinity".

Note also that provided the circle contains the entire cluster there is no size restriction, so we can scale it dynamically, thereby improving the efficiency of the simulation.

3.2.2 On-Lattice: The Fuzzy Annulus

However in the discrete setting, the above approach introduces a bias. As detailed by Loh [14], projecting a circle onto a discrete grid forces the starting positions to snap to the nearest integer coordinates. This creates a 'lattice aliasing' effect (also referred to as a Moireé interference pattern). For large enough clusters, this bias is sufficient to distort key metrics including the fractal dimension.

To eliminate this, we implement Loh's Fuzzy Annulus method. Instead of a fixed radius R, the starting position is sampled from a probability distribution spread over an annulus. This probability distribution utilises the Kaiser-Bessel function I_0, which was developed to suppress high-frequency aliasing artifacts in signal processing [16]:

$$P(x) \propto \frac{I_0(x)}{I_0(R_l)}. \tag{3.1}$$

By sampling from this distribution, we smoothen out the grid aliasing effects, reducing the bias to computationally neglible levels.

Like in the off-lattice case, we can scale the launching annulus dynamically, but to further reduce the bias, the inner boundary of the annulus is at $2R_{max}$.

Remark 3.1 (Rejection Sampling)**.** *Rather than directly sampling from this distribution, we use an optimisation technique called 'rejection sampling', which we explain in appendix A.3.*

3.3 Far-Field Diffusion

Once initialised, or indeed at any point during the diffusion, the walker may be far from the cluster. To improve the efficiency of our algorithms, we exploit the properties of harmonic functinos to quickly return the walker to 'near' the cluster.

3.3.1 Off-Lattice: Re-injection from Infinity

In the continuous case, if the walker strays outside the launching circle ($|z_{out}| > R_l$), we exploit the conformal invariance established in section 2.3 to immediately map it back to this boundary.

We do this by sampling the first hitting point on the circle directly. As shown in corollary 2.3, the return probability is governed by the exterior Poisson kernel. To efficiently sample from this distribution, we first note that the conformal invariance of Brownian motion establishes a statistical equivalence between geometries: to sample the hitting distribution on a physical boundary $\partial\Omega$, we can sample the hitting distribution on a reference domain ∂D (where the

distribution is known, e.g., Uniform) and transform the result via the conformal map $f : D \to \Omega$. This is the key idea for applying conformal maps to our DLA algorithm.

In this instance, we map the uniform exit distribution on the unit circle to the exterior harmonic measure via the following steps:

1. Normalisation: We normalise the walker's position z to $r = |z|/R_l$ and map it to the interior point $a = 1/r$ via inversion. Note that r and a are real numbers.

2. Möbius Transformation: We generate a uniform variate ζ on the unit circle (representing an isotropic exit from the origin). We then apply the real-axis Möbius map:

$$w = \frac{r\zeta + 1}{r + \zeta}. \tag{3.2}$$

   This maps the unit circle to itself, sending the origin to the point a inside the disk. By the conformal invariance of harmonic measure, the transformed variate w is now distributed according to the harmonic measure relative to a.

3. Reconstruction: Inverting now does not change the distribution on the unit circle, so w is distributed exactly according to the exterior harmonic measure. We rotate w by $\arg(z_{out})$ and scale by R_l to obtain the true re-entry position.

This effectively folds the infinite exterior space back onto the launch circle in $\mathcal{O}(1)$ operations.

Remark 3.2. *This 'no-killing' approach was formally introduced by Menshutin and Schur [17] as a 'killing-free' algorithm, where they sampled from the exact hitting distribution (via the Poisson Kernel). We instead adopt Bell's conformal mapping approach, although we map to the unit disk while Bell maps to a straight line.*

3.3.2 On-Lattice: Walk-to-Line

In the discrete setting, we cannot simply use a conformal map to transport a walker from infinity to the cluster's boundary, as the grid lacks the continuous symmetries required for conformal invariance. As an alternative, we employ Loh's Walk-to-Line acceleration method.

When a walker wanders outside the cluster's bounding box, we project its position onto the nearest face of the box. We treat this face as an infinite absorbing line $\partial\Lambda = \{(x, 0) : x \in \mathbb{Z}\}$ and the walker's position as $z_0 = (0, h)$ in the upper half-plane $y > 0$. The problem reduces to sampling the exit probability on the line, the discrete analogue to theorem 2.9.

3.3.2.1 Discrete Method of Images

As established in chapter 2, the exit probability (harmonic measure) at a boundary point $(x, 0)$ is proportional to the discrete flux into that point. Since the potential $u(z)$ vanishes on the boundary, the flux is simply the value of the potential at the immediate interior neighbor:

$$P_{\text{lattice}}(x) \propto \Phi(x, 0) = u(x, 1). \tag{3.3}$$

To calculate this, we must construct a potential $u(z)$ that is harmonic in the upper half-plane

but strictly satisfies the Dirichlet boundary condition $u(x, 0) = 0$ (representing the absorbing line). We achieve this via the Method of Images.

The key idea is to place a virtual "negative walker" (or sink) at the reflected position $(0, -h)$, opposite the true walker at $(0, h)$. Due to the symmetry of the lattice, the potential generated by the real walker at any point on the boundary line $y = 0$ is identical to the potential generated by the image. By subtracting the image potential from the source potential, these contributions cancel out exactly along the line, ensuring $u(x, 0) = 0$ by construction.

Using the discrete potential kernel $a(z)$ defined in section 2.2.3, the resulting field is:

$$u(x, y) = a(x, y + h) - a(x, y - h). \tag{3.4}$$

Substituting this into the flux equation ($y = 1$), we obtain the exact unnormalised probability of the walker hitting the line at relative coordinate x:

$$P_{\text{lattice}}(x) = a(x, 1 + h) - a(x, 1 - h). \tag{3.5}$$

Remark 3.3. *To implement this, we follow Loh exactly. For $h \leq 60$, we evaluate this using his pre-calculated lookup table of $a(z)$. For larger distances, we use the asymptotic expansion of the potential kernel, given in the appendix of [14].*

Remark 3.4. *Computing the normalisation constant for $P_{lattice}(x)$ over the infinite line is computationally problematic. Instead, we once again employ rejection sampling (details in appendix A.3).*

3.3.2.2 The Killing Circle

While the Walk-to-Line method is theoretically valid for any distance, practical implementations are bound by the limits of 64-bit integer arithmetic. If a walker wanders exceptionally far from the cluster (e.g., $R > 10^{14}$ lattice units), the coordinate calculations may overflow.

To prevent numerical instability, we define a hard "Killing Radius" $R_{kill} = 10^{14}$. In the extremely rare event that a walker's random step projects it beyond this limit, the particle is discarded and a new one is initialised. Given the recurrence of 2D random walks, this truncation introduces a vanishingly small statistical error ($\ll 10^{-10}$), which is acceptable for ensuring numerical robustness.

3.4 Near-Field Diffusion

As the walker approaches the cluster, the step size must decrease to avoid overshooting the boundary.

3.4.1 Off-Lattice: Walk-to-Circle

At each step, we calculate the shortest Euclidean distance to the cluster, d_{min}. We can construct a circle C of radius $r = d_{min} - \epsilon$ (where ϵ is a small tolerance). Since the interior of C is empty, the potential is harmonic. By the isotropy of Brownian motion (theorem 2.10 with $z = 0$), the exit distribution is uniform. We therefore advance the walker (at position z) to a random point

on the boundary of C:

$$z_{new} = z + re^{i\theta}, \quad \theta \sim U[0, 2\pi). \tag{3.6}$$

This allows the walker to adaptively skip over large empty regions, taking large steps in the void and small steps in the fjords.

3.4.2 On-Lattice: Walk-to-Square

In the discrete setting, circles do not map cleanly to the grid. Instead, we employ Loh's Walk-to-Square method. This relies on the pre-calculation of exit probabilities for a hierarchy of empty square sizes $L_k = 2^k$ (e.g., from 4×4 up to 512×512).

Unlike a circle, the exit distribution on the boundary of a square is not uniform, as lattice points near the mid-edges are hit more frequently than the corners. We determine these exact probabilities by solving the discrete Dirichlet problem $\mathcal{L}u = 0$ for each square size S_k. For simplicity we used Loh's precomputed solutions (after verification).

Remark 3.5 (Alias Sampling)**.** *To sample from these non-uniform distributions efficiently, we calculate them offline and store them using Walker's Alias Method (detailed in appendix A.4). This allows us to sample the discrete exit location in $\mathcal{O}(1)$ time, regardless of the square size.*

The simulation step proceeds as follows:

1. Determine the largest square S_k centred at the walker that does not intersect the cluster.
2. Query the pre-computed Alias Table for size k to generate a boundary coordinate z_{exit}.
3. Update the walker's position to z_{exit}.

This allows the walker to leap across large voids while preserving the exact non-uniform statistics of the lattice random walk.

3.5 Aggregation

The final stage is the termination of the walk.

3.5.1 On-Lattice: Sticky Boundary

The aggregation condition on $\mathbb{Z}^2$ is straightforward. We define the boundary ∂K as the set of all empty sites adjacent to occupied sites, as we did in chapter chapter 2. The walk terminates when the particle steps onto a site $x \in \partial K$. The site x is then marked as occupied ($x \in K$), and the boundary ∂K is updated locally.

3.5.2 Off-Lattice: Exact Adhering

In the continuous model, aggregation occurs when the walker touches a cluster particle (i.e., the center-to-center distance is 2, since the radius of each particle is 1). Simulating this directly with fixed time-steps is not ideal, due to the potential errors of the walker grazing the cluster or overshooting.

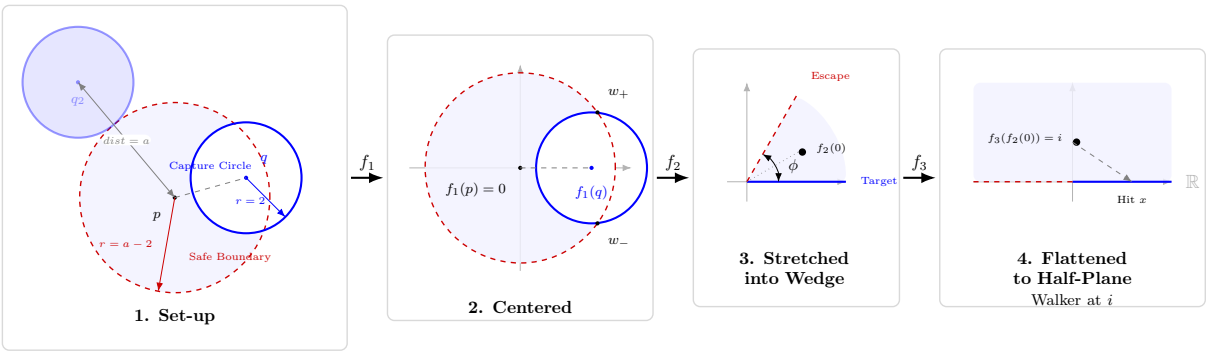

Figure 3.1: The conformal mapping sequence for exact collision. The physical domain (1), formed by the intersection of the safe region and the target exterior, is centred (2), mapped to an infinite wedge (3), and finally opened to the upper half-plane (4) where the exit position is sampled.

To resolve this, we adopt the analytic method described by Bell [15]. When the walker is at position p, close to a target particle centered at q but constrained by a second-nearest neighbour q_2 at distance a, we define a local domain Ω. As illustrated in fig. 3.1 (Stage 1), Ω is the lune formed by the intersection of the "safe escape circle" (radius $a - 2$, centered at p) and the exterior of the "capture circle" (radius 2, centered at q).

To sample the exact exit position from Ω without time-stepping errors, we conformally map this lune to the upper half-plane $\mathbb{H}$ in three stages:

1. **Centering (f_1):** We first translate and rotate the system so that the walker p lies at the origin and the target q lies on the positive real axis:

$$f_1(z) = |q - p| \frac{z - p}{q - p}. \tag{3.7}$$

   As shown in fig. 3.1 (Stage 2), we identify the intersection points of the two circles in this transformed space, denoted w_+ and w_-.

2. **Stretching (f_2):** We apply a Möbius transformation that maps the intersection point w_+ to the origin and w_- to infinity:

$$f_2(z) = \frac{z - f_1(w_+)}{z - f_1(w_-)}. \tag{3.8}$$

   This transformation stretches the circular boundaries into two straight rays emanating from the origin, forming an infinite wedge (fig. 3.1, Stage 3). The angle ϕ between these rays is determined by the intersection geometry. The walker now lies somewhere on the segment of the unit circle within this wedge.

3. **Flattening (f_3):** Finally, we open the wedge into the full upper half-plane using a power map:

$$f_3(z) = z^{\pi/\phi}. \tag{3.9}$$

   This maps the target boundary to the positive real line and the escape boundary to the negative real line (fig. 3.1, Stage 4). Due to the symmetry of the original lune about the $p - q$ axis, the walker is mapped to the point i.

In this final domain, the problem reduces to a standard Brownian motion starting at i and hitting the real line. We sample the hitting point from the Cauchy distribution (derived in section 2.3.4, which easily integrates to give the CDF) and map it back to the physical domain using the inverse composition $f_1^{-1} \circ f_2^{-1} \circ f_3^{-1}$. If the mapped point lies on the target circle, the particle adheres; if it lies on the safe circle, the walker is effectively reset, and diffusion continues.

3.6 The Hybrid Model

To disentangle the source of anisotropy in DLA, we introduce a hybrid model.

3.6.1 Motivation and Design

We hypothesise that the characteristic cross-shape of on-lattice DLA arises primarily from the restriction of particles to fixed grid sites rather than the transport dynamics. To test this, we develop a model where the transport is purely isotropic (Brownian motion), but the final aggregation is constrained to a grid geometry.

The critical design choice is the aggregation mechanism. We evaluated three potential strategies:

1. The Sticky Boundary: A walker aggregates if it enters any square cell within the sticky boundary. We rejected this because the square capture zone imposes lattice anisotropy onto the incoming probability flux before aggregation occurs, preventing a clean separation of transport and packing effects.

2. The Antennae Method: As proposed by Menshutin and Shchur [10], particles only stick if they hit a lattice site along specific axes; otherwise, they are reflected. This is similar in some sense to "Noise-Reduced DLA" [18] which was designed to accelerate convergence to asymptotic forms. We don't use this method as it potentially alters both the rate of convergence of the cluster (as in Noise-Reduced DLA) and the harmonic measure itself, since walkers are artificially redirected.

3. Rotate and Snap **(Chosen)**: The walker behaves exactly as a hard sphere in continuous space. Upon collision, the particle relaxes into the nearest valid lattice site. This method preserves the exact continuous harmonic measure of the cluster until the instant of impact, ensuring the arrival statistics are purely isotropic.

Remark 3.6. *One might question whether the specific choice of aggregation rule significantly alters the growth. Theoretical estimates suggest that these microscopic differences are likely negligible.*

As derived from Beurling's estimates on harmonic measure [19], once a random walker enters the immediate vicinity of a boundary point, the probability of it escaping to hit a different, distant site decays rapidly. Specifically, if a walker is at a small distance δ from a cluster site P, the harmonic probability ω that it diffuses away to hit a different site Q (at distance $r \gg \delta$) is bounded by:

$$\omega(z \to Q) \lesssim \left(\frac{\delta}{r}\right)^{\alpha} \tag{3.10}$$

where $\alpha > 0$ depends on the local geometry.

This implies that the decision of where to aggregate is effectively determined by the macroscopic flux before the walker resolves the microscopic shape of the target site. Therefore the exact adherence rule is likely not overly important, although we still prefer our method for the reasons outlined previously.

3.6.2 Implementation

The algorithm proceeds identically to the Off-Lattice simulation until the collision condition $|z_{walker} - z_{cluster}| \leq 2$ is met (since particle radius $r_p = 1$).

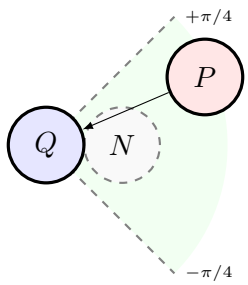

Figure 3.2: The angular window for hybrid aggregation. If a walker P collides with particle Q within the angle $|\phi| < \pi/4$ (shaded region), it is snapped to the nearest lattice slot N. We can assume N is empty because if it were occupied, its own collision shell would have intercepted the walker before it could reach Q within this specific angle.

Let the continuous collision angle relative to the target particle center be $\phi \in [0, 2\pi)$. We project this continuous angle onto the discrete set of lattice directions $\Theta = \{0, \pi/2, \pi, 3\pi/2\}$ by minimizing the angular distance:

$$\theta_{snap} = \frac{\pi}{2} \left\lfloor \frac{4\phi + \pi}{2\pi} \right\rfloor . \tag{3.11}$$

The particle is then rigidly placed at the lattice coordinate:

$$z_{new} = z_{cluster} + 2e^{i\theta_{snap}}. \tag{3.12}$$

A potential concern is whether the "snapped" lattice site is available. We provide a simple geometric argument to justify that the target slot is always empty.

Consider a walker P hitting a cluster particle Q at angle ϕ. The snapped destination is the lattice neighbour N closest to the point of impact. For a grid of spacing 2, the neighbour N is located at distance 2 from Q (centre-to-centre). If N were already occupied, its own collision shell (radius 2) would overlap the collision shell of P.

Simple trigonometry shows that for any impact angle ϕ closer to N than any other neighbour (i.e., $|\phi| < \pi/4$), the distance from the impact point to N is strictly less than 2:

$$|z_{hit} - z_N| = |2e^{i\phi} - 2| = 2\sqrt{2 - 2\cos\phi}. \tag{3.13}$$

For $|\phi| < \pi/4$, this distance is at most $2\sqrt{2 - \sqrt{2}} \approx 1.53$, which is well inside the collision radius of 2. This implies that if N were occupied, the walker would have intersected N's collision shell

before reaching P's surface. Since the algorithm halts at the first contact, the target slot N must be empty.

Chapter 4

Simulation Results

Chapter 5

Fractal Dimension of DLA

5.1 Introduction to Dimension

To understand the fractal dimension of a Diffusion-Limited Aggregation (DLA) cluster, we must first establish a mathematically rigorous definition of "dimension." Intuitively, dimension describes how an object fills space as it scales. For standard Euclidean geometries, this scaling relationship resolves to an integer. If we scale a 2D solid square by a factor of 2, its area (or "mass") increases by $2^2 = 4$. For fractal objects, this scaling relationship does not resolve to an integer, leading to a fractional dimension D where the mass scales by 2^D. To formally quantify this, we rely on the framework established by Felix Hausdorff and later refined for practical application by Minkowski and Bouligand.

5.1.1 Hausdorff Dimension

The formal mathematical foundation for fractal dimension is the Hausdorff dimension, which relies on the concept of optimal covers. Following Falconer [20], we first define the Hausdorff measure.

Definition 5.1 (Hausdorff Measure)**.** *Let $F \subset \mathbb{R}^n$ and $s \geq 0$. A countable collection of sets $\{U_i\}$ is a δ-cover of F if $F \subset \bigcup_{i=1}^{\infty} U_i$ and the diameter $|U_i| \leq \delta$ for all i. The s-dimensional Hausdorff measure is defined as:*

$$\mathcal{H}^s(F) = \lim_{\delta \to 0} \inf \left\{ \sum_{i=1}^{\infty} |U_i|^s : \{U_i\} \text{ is a } \delta\text{-cover of } F \right\} \tag{5.1}$$

This definition asks us to cover the fractal with sets (like tiny circles or blobs) of maximum size δ, summing the s-th powers of their diameters. The infimum ensures we are finding the absolute most efficient possible cover.

Falconer demonstrates that as s varies, $\mathcal{H}^s(F)$ exhibits a critical "jump" behavior, which fundamentally defines the dimension of the set.

Proposition 5.1 (Hausdorff Dimension)**.** *There exists a unique critical value, called the Hausdorff*

dimension $\dim_H(F)$*, such that for any set* F*:*

$$\mathcal{H}^s(F) = \infty \quad \textit{if } s < \dim_H(F)$$
$$\mathcal{H}^s(F) = 0 \quad \textit{if } s > \dim_H(F)$$

Formally, $\dim_H(F) = \inf\{s \geq 0 : \mathcal{H}^s(F) = 0\}$.

While mathematically beautiful and rigorously applicable to any set, calculating $\dim_H(F)$ requires finding the optimal variable-sized cover as $\delta \to 0$. This optimization is computationally intractable for complex physical systems like discrete DLA simulations.

5.1.2 Box-Counting Dimension (D_{box})

To extract a measurable metric from these abstract concepts, we turn to the Box-Counting (or Minkowski-Bouligand) dimension. Rather than starting with a formal limit, it is more intuitive to begin with the fundamental power law of scaling.

Consider covering a standard Euclidean object with a uniform grid of boxes of side length ϵ. If the object is a 1D line of length L, the number of boxes required is $N(\epsilon) \propto L/\epsilon \sim \epsilon^{-1}$. If the object is a 2D square of area A, the number of boxes is $N(\epsilon) \propto A/\epsilon^2 \sim \epsilon^{-2}$. We can generalize this scaling relationship for an object of arbitrary dimension D_{box} as:

$$N_\epsilon(F) \sim \epsilon^{-D_{box}} \tag{5.2}$$

To isolate the dimension D_{box}, we take the logarithm of both sides, yielding $\log N_\epsilon(F) \approx -D_{box} \log \epsilon + c$, where c is a constant. Rearranging for D_{box} gives:

$$D_{box} \approx \frac{\log N_\epsilon(F) - c}{-\log \epsilon} \tag{5.3}$$

As the box size becomes infinitesimally small ($\epsilon \to 0$), the constant c becomes negligible compared to the logarithm of ϵ. This intuitive derivation directly motivates the formal mathematical definition outlined by Falconer [20]:

Definition 5.2 (Box-Counting Dimension)**.** *Let* $N_\epsilon(F)$ *be the smallest number of* ϵ*-mesh cubes (boxes of side length* ϵ*) that intersect a set* F*. The Box-Counting dimension* D_{box} *is defined as:*

$$D_{box} = \lim_{\epsilon \to 0} \frac{\log N_\epsilon(F)}{-\log \epsilon} \tag{5.4}$$

assuming this limit exists.

Practically, by measuring the number of occupied boxes $N_\epsilon(F)$ for various box sizes ϵ and plotting them on a log-log graph, D_{box} is extracted directly from the negative slope of the regression line.

5.1.3 Growth Rate Dimension (D_β)

While Box-Counting analyzes the "frozen" final geometry of a cluster by varying the measurement resolution, most modern literature on DLA utilizes a kinetic proxy for dimension based on the

cluster's *growth rate* over time.

To quantify this dynamic expansion, we must first define a robust metric for the spatial extent of the cluster. While the maximal radius (the absolute distance to the furthest particle) is highly susceptible to stochastic outliers and noisy tip-growth, the radius of gyration provides a stable, averaged measure of the mass distribution.

Definition 5.3 (Radius of Gyration)**.** *For a discrete aggregate of N identical particles located at positions $\mathbf{r}_i$, let $\mathbf{r}_{cm} = \frac{1}{N}\sum_{i=1}^{N}\mathbf{r}_i$ be the center of mass. The radius of gyration, R_g, is the root-mean-square distance of the particles from the center of mass, defined as:*

$$R_g = \sqrt{\frac{1}{N}\sum_{i=1}^{N}|\mathbf{r}_i - \mathbf{r}_{cm}|^2} \tag{5.5}$$

As a cluster aggregates, we can continuously characterize its overall spatial extent using this radius of gyration, R_g, and measure its expansion as a function of the total number of aggregated particles, N. This dynamic expansion is described by the scaling relation:

$$R_g \sim N^{\beta} \tag{5.6}$$

where β is defined as the growth rate exponent.

To connect this kinetic growth rate back to the geometric power law established in the previous section, consider the mass-radius relationship of a fractal. In the Box-Counting framework, the mass (number of particles N) contained within a characteristic radius R scales as $N \sim R^{D_{box}}$. Rearranging this spatial scaling to solve for the radius gives:

$$R \sim N^{1/D_{box}} \tag{5.7}$$

By equating our dynamic radial expansion (R_g) with the geometric spatial radius (R), we observe that the kinetic growth exponent relates to the dimension via $\beta = 1/D_{box}$. This powerful mathematical link justifies the definition of a distinct Growth Rate Dimension, D_β, defined simply as the inverse of the growth rate:

$$D_\beta = \frac{1}{\beta} \tag{5.8}$$

5.1.4 The Link Between D_{box} and D_β

Because of the algebraic link established above ($1/\beta = D_{box}$), it is a common assumption in the DLA literature that the geometric box-counting dimension and the kinetic growth dimension represent identical concepts and yield equivalent values ($D_{box} = D_\beta$). For an ideal, perfectly isotropic fractal, these two metrics do indeed converge to the known off-lattice DLA dimension of $D \approx 1.71$.

However, similar underlying definitions do not guarantee identical physical concepts. As highlighted by Grebenkov and Beliaev [8], the two metrics view the cluster through fundamentally different lenses: D_{box} measures the bulk space-filling capacity of the interior, while D_β measures the maximal radial expansion at the kinetic frontier.

They provide a simple but illuminating counter-example to prove that these dimensions can

entirely decouple for anisotropic structures. Consider an aggregate composed of $2t$ particles. Half of these particles (t) form a dense 2D disk of radius $R_{disk} \propto \sqrt{t}$. The other half ($t$) form a 1D interval (a line) of length $L \propto t$.

- **Box-Counting (D_{box}):** The overall fractal dimension is determined by the densest, most space-filling component. Because the aggregate contains a solid 2D disk, its Box-Counting dimension is $D_{box} = 2$.
- **Growth Rate (D_β):** The growth rate is determined by the maximal spatial extent. The 1D interval length ($L \propto t$) grows radially much faster than the disk ($R_{disk} \propto \sqrt{t}$). Therefore, the overall radius scales as $R \propto t$. Since the total mass is $N = 2t \propto t$, the radius scales linearly with the mass ($R \propto N^1$). Thus, the growth rate is $\beta = 1$, yielding a Growth Rate Dimension of $D_\beta = 1/\beta = 1$.

In this aggregate, $D_\beta = 1$ whilst $D_{box} = 2$. As Grebenkov and Beliaev note, a naive rescaling of this object by its maximal diameter ($\propto t$) results in a trivial limit (the unit interval) because "the part with the higher dimension but smaller growth rate (the disk) is shrunk and thus fully eliminated in the limit $t \to \infty$" [8].

While DLA clusters are vastly more complex than a simple line-and-disk shape, this exact phenomenon—long, fast-growing arms dictating the radial expansion rate (β) while a denser interior structure dictates the bulk mass scaling (D_{box})—provides the theoretical foundation for why a significant discrepancy between these global metrics might emerge as clusters evolve anisotropically.

5.2 Methodology

To compute both D_{box} and D_β, we rely on extracting slopes from log-log relationships. However, raw data from discrete DLA simulations contains inherent structural artifacts that can heavily bias standard regression models. To accurately capture the true scale-invariant fractal dimension, stringent methodological constraints and robust statistical techniques must be applied.

5.2.1 Calculating Box Dimension: The Box Size Problem

The Box-Counting dimension is calculated by generating grids of varying box sizes (ϵ), counting the occupied boxes $N(\epsilon)$, and plotting $\log N(\epsilon)$ against $\log \epsilon$. A primary point of contention in the literature is exactly which box sizes to include in this regression.

One cannot naively use all mathematically possible box sizes. As noted by Grebenkov and Beliaev [8], DLA clusters are bounded by finite-size effects at the macroscopic scale and by lattice discretization at the microscopic scale. If ϵ is too large (approaching the size of the cluster), $N(\epsilon) \to 1$, which trivially yields a slope of 0. Conversely, if ϵ is too small (approaching the size of a single particle or lattice site), the grid merely counts individual discrete pixels, resolving to the Euclidean dimension of the particles themselves rather than the fractal dimension of the aggregate.

To systematically identify the valid scaling regime, we compare two different regression techniques: standard Ordinary Least Squares (L2) regression and the robust Theil-Sen estimator.

While L2 minimizes the sum of squared residuals, making it highly sensitive to outliers, the Theil-Sen estimator calculates the median of all pairwise slopes between data points, rendering it exceptionally robust to boundary artifacts.

When performing regressions that include small box sizes ($\epsilon < 16$ pixels), a consistent and significant deviation emerges between the two techniques. The L2 regression slope is consistently pulled down (yielding an artificially inflated dimension) by the dense, space-filling nature of individual discrete particles at the smallest scales. The Theil-Sen regression resists this localized pull. This systematic discrepancy provides empirical proof that microscopic lattice effects dominate at $\epsilon < 16$. Consequently, our methodology strictly imposes a lower bound of $\epsilon \geq 16$ pixels for all D_{box} calculations, ensuring the extracted slope reflects the true bulk fractal geometry rather than microscopic "pixel noise."

Figure Placeholder: Box-Counting Regression Comparison
Insert Notebook Plot: Log-Log scatter of $N(\epsilon)$ vs ϵ for a single cluster, showing L2 vs Theil-Sen fits. Highlight the deviation at $\epsilon < 16$.

Figure 5.1: Comparison of L2 and Theil-Sen regressions for Box-Counting dimension on a single cluster. The divergence at small box sizes ($\epsilon < 16$) illustrates the susceptibility of L2 to microscopic lattice artifacts, justifying the choice of a robust estimator and lower bound cutoffs.

5.2.2 Calculating Growth Rate Dimension: Sampling and Uncertainty

For the kinetic evaluation (D_β), the radius of gyration R_g is tracked as a function of the aggregate mass N. Because simulating up to $N = 10^7$ generates an immense amount of discrete data points, the series is logarithmically sampled to maintain a computationally efficient and structurally representative dataset. By extracting 50 samples per logarithmic decade of N, a cluster growing to 10^7 particles yields a manageable set of approximately 350 data points upon which to perform the regression.

Even with this sampling, the stochastic nature of tip-growth causes random fluctuations in R_g that can influence standard L2 fitting. By calculating both the Theil-Sen and L2 regressions on this sampled data, we not only extract a robust median slope via Theil-Sen, but we also gain a powerful diagnostic tool. The absolute difference between the Theil-Sen and L2 slopes, $\Delta D = |D_{Theil-Sen} - D_{L2}|$, serves as a non-parametric measure of structural uncertainty. A large discrepancy indicates that the object lacks perfect scale-invariance and possesses a high degree of internal structural variation or noise. This difference allows us to rigorously define the error

bounds of our dimension metrics before analyzing the global results in the subsequent section.

5.3 Results and Statistical Errors

Table 5.1: Summary of final dimensional metrics for 10^6 (1M) and 10^7 (10M) particle clusters across different simulation environments. Values are presented as the ensemble mean $\pm$ one standard deviation.

Model (Size)	**Growth Rate (β)**	**Growth Dim (D_β)**	**Box Dim (D_{box})**
Off-Lattice (1M)	0.582 ± 0.008	1.717 ± 0.025	1.656 ± 0.010
Off-Lattice (10M)	0.583 ± 0.007	1.717 ± 0.022	1.666 ± 0.013
Lattice (1M)	0.587 ± 0.009	1.703 ± 0.027	1.661 ± 0.011
Lattice (10M)	0.590 ± 0.008	1.696 ± 0.024	1.670 ± 0.013
Hybrid (1M)	0.589 ± 0.009	1.698 ± 0.027	1.661 ± 0.011
Hybrid (10M)	0.592 ± 0.008	1.689 ± 0.023	1.670 ± 0.013

The results demonstrate that the Off-Lattice control model remains highly stable and isotropic, yielding $D_\beta = 1.717 \pm 0.022$ at $N = 10^7$, matching theoretical expectations. In stark contrast, both the Lattice and Hybrid models exhibit significant "drift," where the dimension systematically decreases as the cluster grows.

The histogram distributions (Figure 5.2) visually confirm a critical physical mechanism: the Hybrid model's distribution almost perfectly overlays the Lattice model's distribution. This empirically proves that the underlying random walk continuous nature is largely irrelevant to the final fractal dimension; the rigid geometry of the sticking rules alone dictates the dimensional drift.

5.4 Comparing the Evolution of Dimension and Growth Rate

By analysing the dimension at logarithmically spaced snapshots during the clusters' growth, we can observe the temporal dynamics of the spatial structure. A key finding of this study is the apparent "lag" between the kinetic dimension (D_β) and the bulk geometric dimension (D_{box}) in anisotropic models.

As shown in Figure 5.3, while the two metrics track closely during early growth ($N < 10^5$), they exhibit distinct behaviors at late stages. The growth rate dimension, D_β, drops sharply. Because D_β is governed by the expanding radius of gyration, it is immediately sensitive to the accelerated growth of the distinct lattice "arms" piercing outward. Consequently, D_β also exhibits a higher variance, as it is driven by the highly stochastic, noisy nature of tip-growth.

Conversely, D_{box} relies on the summation of the entire frozen structure. While the tips are accelerating outward, the deep interior fjords of the cluster continue to slowly fill with particles that bypass the outer screening. Therefore, the geometric box dimension lags behind the kinetic measurement, updating much slower to the structural changes and exhibiting significantly lower variance across the ensemble. This discrepancy highlights that global dimension metrics ultimately fail to capture the localized axial acceleration occurring at the lattice boundaries—a phenomenon that necessitates a direct quantification of angular anisotropy, which is explored in the following

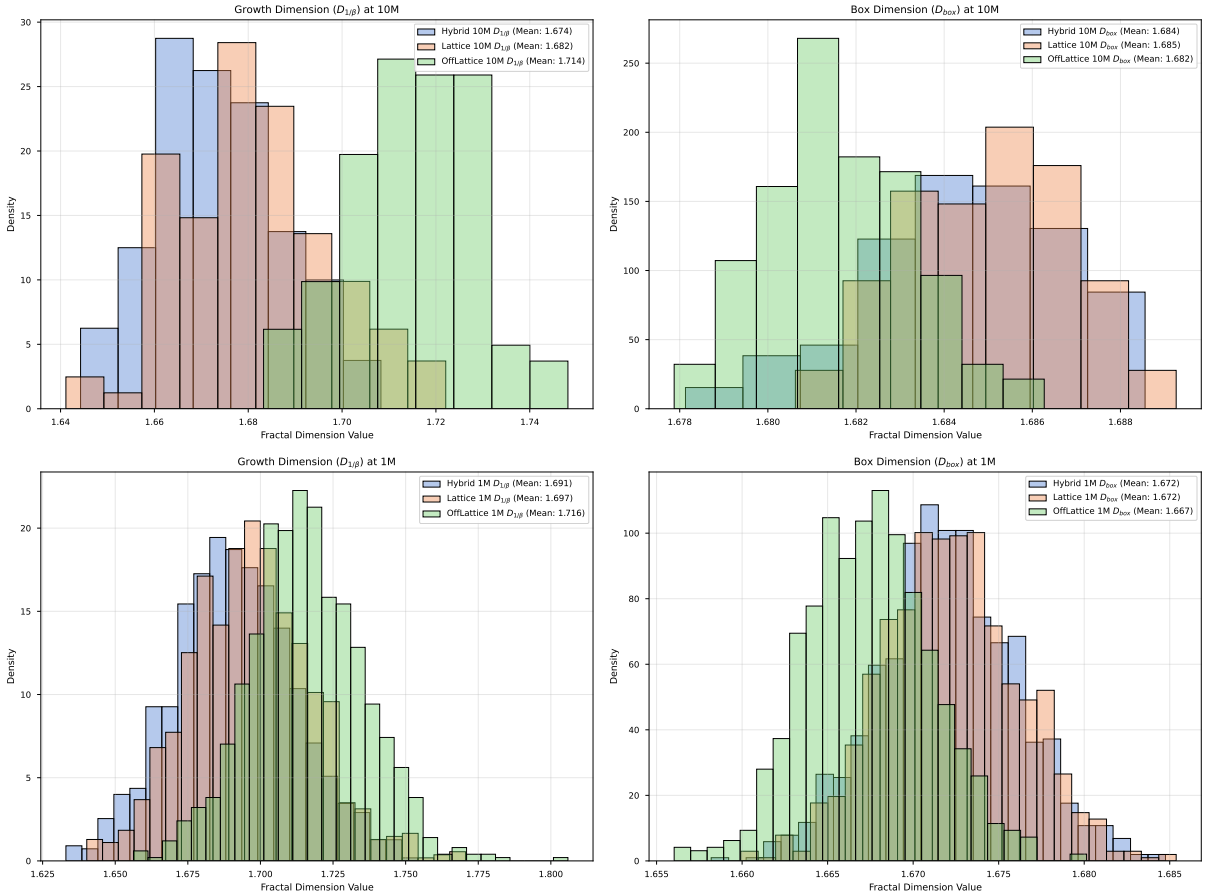

Figure 5.2: Probability density distribution of the Growth Dimension (D_β) and Box-Counting Dimension (D_{box}) for 10^6 (1M) and 10^7 (10M) clusters. The Hybrid model closely mirrors the Lattice model, diverging significantly from the Off-Lattice control in both metrics and across both sizes.

Figure Placeholder: Dimension Evolution Over Time

Insert Notebook Plot: plot_batch_dimension_evolution() showing the temporal lag.

Figure 5.3: Statistical evolution of D_{box} and D_β against cluster size N for the Lattice model. Shaded regions represent $\pm 1\sigma$.

chapter.

Chapter 6

Anisotropy in DLA

6.1 What is Anisotropy?

6.2 Angular Radius of Gyration

6.3 Angular Growth Rate Analysis

6.4 Conclusions

Bibliography

[1] T. A. Witten and L. M. Sander. Diffusion-limited aggregation, a kinetic critical phenomenon. *Physical Review Letters*, 47(19):1400, 1981.

[2] Mitsugu Matsushita et al. Fractal structures of zinc metal leaves grown by electrodeposition. *Phys. Rev. Lett.*, 53:286–289, 1984.

[3] Eshel Ben-Jacob et al. Generic modelling of cooperative growth patterns in bacterial colonies. *Nature*, 368:46–49, 1994.

[4] Lincoln Paterson. Diffusion-limited aggregation and the Saffman-Taylor instability. *Phys. Rev. Lett.*, 52:1621–1624, 1984.

[5] Gregory F. Lawler and Vlada Limic. *Random Walk: A Modern Introduction.* Cambridge University Press, 2010.

[6] Paul Meakin. Diffusion-controlled cluster formation in 2-6-dimensional space. *Phys. Rev. A*, 27:1495–1507, 1983.

[7] Harry Kesten. How long are the arms in DLA? *J. Phys. A: Math. Gen.*, 20:L29–L33, 1987.

[8] D. S. Grebenkov and D. Beliaev. How anisotropy beats fractality in two-dimensional on-lattice diffusion-limited-aggregation growth. *Phys. Rev. E*, 96:042159, 2017.

[9] A. S. Shmyrov, A. I. Shmyrova, and E. S. Gerasimov. Hybrid L-system–DLA scheme for modeling morphology of fungal colonies. *Physica A: Statistical Mechanics and its Applications*, 517:443–452, 2019.

[10] Anton Yu Menshutin and Lev N. Shchur. Morphological diagram of diffusion driven aggregate growth in plane. *Comp. Phys. Comm.*, 182:1819–1821, 2011.

[11] David Williams. *Probability with Martingales.* Cambridge University Press, Cambridge, 1991.

[12] Lars V. Ahlfors. *Complex Analysis: An Introduction to the Theory of Analytic Functions of One Complex Variable.* McGraw-Hill, New York, 3rd edition, 1979.

[13] Peter Mörters and Yuval Peres. *Brownian Motion.* Cambridge University Press, Cambridge, 2010.

[14] Yen Lee Loh. Bias-free simulation of diffusion-limited aggregation on a square lattice. *Computer Physics Communications*, 185(3):841–846, 2014.

[15] James Henry Bell. *Simulation of Diffusion Limited Aggregation Models and Related Results.* Phd thesis, University of Cambridge, October 2019.

[16] James F Kaiser. Digital filters. In F.F. Kuo and J.F. Kaiser, editors, *System Analysis by Digital Computer*, pages 218–285. Wiley New York, 1966.

[17] Anton Yu Menshutin and Lev N Shchur. Test of multiscaling in a diffusion-limited-aggregation model using an off-lattice killing-free algorithm. *Physical Review E*, 73(1):011407, 2006. doi: 10.1103/PhysRevE.73.011407.

[18] Paul Meakin. Noise-reduced diffusion-limited aggregation. *Physical Review A*, 36(1):332, 1987.

[19] Harry Kesten. Hitting probabilities of random walks on $\mathbb{Z}^d$. *Stochastic Processes and their Applications*, 25:165–184, 1987.

[20] Kenneth Falconer. *Fractal Geometry: Mathematical Foundations and Applications.* John Wiley & Sons, Chichester, England, 1990.

[21] Milton Abramowitz and Irene A. Stegun, editors. *Handbook of Mathematical Functions: With Formulas, Graphs, and Mathematical Tables.* Dover Publications, New York, 1965. Eq. 9.8.1.

Appendix A

Implementation Details and Optimisations

While chapter 3 outlines the mathematical principles of the algorithms, the practical implementation requires specific data structures and optimisations to achieve the necessary scale. This appendix details the computational architecture of the simulation code, written in Python and accelerated using the Numba JIT (Just-In-Time) compiler.

The lattice and off-lattice versions of the algorithm are versions of code created by Loh [14] and Bell [15] respectively. However, these were originally done in C and C++, and in translating to Python our implementation diverges both to account for differences in the language structure, and to account for the limitations of our computer (a 2017 Macbook with 8GB RAM). We will highlight some of these changes here, as well as explain some of the more programming-related concepts and structures behind the algorithms.

The full source code for all models and analysis scripts is available at `https://github.com/RamsayDavis/dla`.

A.1 Computational Framework

A.1.1 The Numba JIT Compiler

The simulations are implemented in Python to leverage its flexibility, to consistent with the analysis section of this project, and as it is the language we are most comfortable with. The core numerical routines are decorated with `@njit(fastmath=True)`. This compiles the Python bytecode into optimised LLVM machine code at runtime, allowing for performance parity with C++ while maintaining a modular, functional codebase. Our code is likely somewhat slower than the original implementations, but this mitigates this difference, and allows us to generate clusters of size 10^7 particles in about 10 to 15 minutes on our CPU.

A.1.2 Rounding Semantics

One minor but important detail to consider in stochasti simulations is the rounding behaviour of floating-point co-ordinates, particularly in the lattice and hybrid models where the output is in integers.

- Python: Python 3 employs "Banker's Rounding" (round half to even). For example, `round(2.5) -> 2` and `round(3.5) -> 4`. This introduces a subtle statistical bias in random walks, as even coordinates are slightly favoured over odd ones.
- C and C++: The reference C++ implementations use the standard library `round()` function, which rounds half away from zero (e.g., `round(2.5) -> 3`).

To ensure our random walks match the theoretical lattice statistics and the reference datasets exactly, we implemented a custom kernel `_round_cpp(x)` that explicitly emulates the C++ behaviour in Python.

A.2 Spatial Indexing and Hierarchical Skipping

In the diffusion steps of all versions of DLA, we require the distance from the walker to the nearest particle in the cluster, in order to work out the maximum allowed step-size. Additionally it is needed in the off-lattice algorithm for the adhesion step, since we need the two nearest particles.

To do this efficiently, we employ a Hierarchical Map. The exact implementation differs between the two models, but the essential idea is we overlay different sized grids over the cluster and surrounding area, and record whether the cells in each grid are occupied or not. By traversing the different grid layers, we can quickly work out these distances.

A.2.1 Lattice: Flattened Bit-Map

Data Structure: In the discrete model, the hierarchy must have enough layers to view individual lattice sites at the finest level. Loh's code uses a 2D array of pointers, one per hierarchy level. However for large grids this can be problematic in Python, with memory fragmentation and cache locality issues.

We instead implement a Flattened 1D Bit-Map. All levels are packed into a single contiguous `uint8` array, `flat_ss`, indexed via an offset table. We track occupancy at the bit level (1 bit per site).

Linear Search: We utilise pre-computed Green's function tables for jumps of radius 2^0 to 2^8 (corresponding to the squares of sizes 2 to 512). The routine `_diffuse_variable_flat` performs a linear scan starting from the coarsest useful level (pixel size ≈ 256) down to the finest:

1. Start: Check level $L_{start} = \max(2, L_{max} - 9)$.
2. Check: If the 3×3 neighbourhood is empty, the walker is in a void of radius ≥ 256. We immediately apply the maximum jump ($R = 256$) and exit.
3. Repeat: If occupied, we increment the level index to check finer resolutions ($R = 128, 64 \ldots 1$).

While L_{max} is often greater than 9, we only have precalculated tables for up to this size, hence the restriction. A potential improvement to the on-lattice code would be to include pre-calculated tables of larger sizes, but this was not a priority for achieving our goal cluster sizes.

A linear scan is preferred over binary search here because DLA walkers spend most time in voids, where the first (coarsest) check succeeds immediately.

A.2.2 Off-Lattice: Dual-Structure Strategy

The continuous model separates "Skipping" and "Hitting" into two distinct structures to optimise for dynamic memory usage.

A.2.2.1 1. The Coarse Hierarchy (Skipping)

Structure: A dynamic list of 2D arrays. Unlike the lattice version, we do not flatten this into a single array. At the finest level, the mesh width is approximately 24 units, meaning the memory pressure is orders of magnitude lower than the unit-resolution lattice grid.

Binary Search: Since continuous walkers can jump any distance, we utilise the full depth of the hierarchy. We use a binary search to efficiently locate the transition layer between "Occupied" and "Empty". This guarantees $\mathcal{O}(\log L)$ performance, ensuring stability whether the walker is in a void or a deep fjord. The greater number of layers makes a binary search worthwhile here.

A.2.3 2. The Spatial Hash (Hitting)

Structure: When the hierarchy indicates the walker is close to the cluster, we require the exact particle coordinates for the adhesion step. However, the coarse hierarchy only records occupancy and not the precise coordinates. To query nearest neighbours in $\mathcal{O}(1)$ time without dynamic memory allocation, we implement a static chained hash table on a fine grid.

This mimics a “List of Lists" using two pre-allocated integer arrays:

1. `grid_head[cell_index]`: Stores the index of the most recently added particle in the grid cell. Initialised to -1.
2. `particle_next[particle_index]`: Stores the index of the next particle in the same cell, forming a linked list.

Insertion: When a new particle P_{new} is added to cell C:

$$\texttt{particle_next}[P_{new}] = \texttt{grid_head}[C] \quad ; \quad \texttt{grid_head}[C] = P_{new} \tag{A.1}$$

Bell’s code used C++ vectors, with one vector per cell in the finest layer of the hierarchy. We prefer our linked list approach as we avoid resizing arrays or allocating objects. In both cases we can iterate through the (linked) list to find the closest particle to the walker.

A.3 Rejection Sampling

A.3.1 Concept and Motivation

Several probability distributions required for the on-lattice simulation possess Cumulative Distribution Functions (CDFs) that are computationally expensive or impossible to invert analytically.

A naive approach would be to discretise the domain and build a lookup table (numerical integration). However, this introduces systematic discretisation error. Instead, we employ rejection sampling. This method allows us to generate samples from a complex target distribution $f(x)$ using a simpler "proposal" distribution $g(x)$, producing exact samples from $f(x)$ without any approximation error.

A.3.2 The Algorithm

Given a target density $f(x)$ and an envelope density $g(x)$ such that $f(x) \leq Mg(x)$ for all x (where M is a constant scaling factor), the algorithm proceeds as follows:

Algorithm 1: Rejection Sampling

```
Function SampleRejection(f, g, M):
    while True do
        X ← Sample(g) ;                          // Propose from envelope
        U ← RandomFloat(0, 1) ;                  // Verification variable
        if U ≤ f(X) / (M·g(X)) then
            return X ;                           // Accept and return
        // Else: Reject and loop again
```

Geometrically, this is equivalent to generating uniform random points under the curve $Mg(x)$ and keeping only those that fall under the curve $f(x)$. While the runtime is non-deterministic (dependent on the acceptance rate $1/M$), it is often computationally superior to numerical inversion because it avoids evaluating expensive special functions (like integrals of Bessel functions) for every rejected sample.

A.3.3 Application 1: The Fuzzy Annulus

To eliminate lattice aliasing, we launch particles from a radial distribution proportional to the Kaiser-Bessel function I_0.

- $f(x)$**:** $I_0(\beta\sqrt{1-x^2})$.
- $g(x)$**:** A Gaussian distribution $\mathcal{N}(0, \sigma^2)$ that strictly bounds I_0.

We sample candidate radii from the Gaussian (using standard library functions) and reject them based on the ratio of the Bessel weight to the Gaussian weight.

A.3.4 Application 2: Walk-to-Line

When a walker projects onto the bounding box face, we must sample the exit position on the line $y = 0$.

- $f(x)$**:** The exact discrete flux calculated via the Method of Images ($P_{lattice} \propto u(x, 1) - u(x, -1)$).
- $g(x)$**:** The Cauchy distribution (the analytical solution for continuous Brownian motion).

As established by Loh [14], the continuous flux serves as a valid envelope for the discrete flux. By setting the scaling constant $M \approx 1$, we ensure the condition $f(x) \leq Mg(x)$ holds, resulting in a computationally efficient sampling routine with a high acceptance rate.

A.4 Walker's Alias Method

A.4.1 Motivation

For the "Walk-to-Square" algorithm, we must sample the exit location of a walker from the boundary of a square. This distribution is non-uniform (the exit probability is highest at the midpoints of the edges and lowest at the corners).

A naive approach be to search a Cumulative Distribution Function (CDF). For a square of perimeter N, this requires a binary search, scaling as $\mathcal{O}(\log N)$. As the square size increases (e.g., $N = 2048$ for a 512×512 square), this $\log N$ cost inside the inner loop becomes a significant bottleneck.

To eliminate this overhead, we follow Loh and implement Walker's Alias Method, which reduces the sampling complexity to $\mathcal{O}(1)$—constant time, regardless of the square size.

Essentially the Alias Method works by reshaping the uneven probability histogram into a perfectly rectangular block.

Imagine the probability mass function as a set of vertical bars. We normalize them so the average height is 1. The method then "slices" the excess height from "over-full" bins (where $P > 1$) and moves these blocks to fill the gaps in the "under-full" bins (where $P < 1$).

The result is a table of size N where every bucket contains exactly the same total probability mass ($1/N$), but that mass is shared between at most two outcomes: the original occupant of that bucket, and the overflow from one other bin.

A.4.2 Implementation

We pre-compute two lookup tables for every square size $L = 2^k$:

- **Prob Table (U_i):** The threshold probability for keeping the original index i.
- **Alias Table (K_i):** The alternate index if the threshold check fails.

Sampling is then performed by first a discrete uniform random sample, and then a continuous uniform sample which is compared to the threshold value:

Algorithm 2: Alias Sampling (Constant Time)

```
Function SampleAlias(U, K):
    i ← RandomInteger(0, N − 1) ;          // Select a bucket uniformly
    y ← RandomFloat(0, 1) ;                // Select height within bucket
    if y < U[i] then
        return i ;                         // Return bucket owner
    else
        return K[i] ;                      // Return bucket alias
```

A.5 Numerical Approximations

A.5.1 Custom Bessel Functions

The Kaiser-Bessel launching distribution requires the Modified Bessel Function of the first kind, $I_0(x)$. The standard Python implementation, `scipy.special.i0`, is not supported in Numba's `nopython` mode (as it links to C Fortran libraries not visible to the JIT compiler). We therefore implemented the polynomial approximation from Abramowitz and Stegun [21] (Eq. 9.8.1) directly in Python/Numba. This ensures that the launching bias correction is applied efficiently without breaking the JIT compilation chain or requiring callbacks to the Python interpreter.

A.6 Usability and Data Management

To facilitate large-scale parameter sweeps and ensure consistency between the discrete and continuous datasets, we introduced two workflow improvements to the reference implementation.

A.6.1 Heuristic Grid Sizing

The spatial extent of a DLA cluster is stochastic and difficult to predict a priori. Loh's code requires the user to manually define the grid dimension L_{max}. Underestimating this parameter leads to runtime failures (truncation), while overestimating it increases memory overhead.

To automate this, we implemented a heuristic initialisation step. Before memory allocation, we estimate the expected cluster radius R_{est} using the known fractal dimension $D_f \approx 1.71$ and apply a safety factor of 1.5 to account for the irregularity of the dendritic arms:

$$L_{max} = \lceil \log_2(1.5 \times 2.2 \times N^{1/1.71}) \rceil \tag{A.2}$$

This ensures the simulation domain dynamically scales to accommodate the target particle count N without manual intervention.

A.6.2 Standardised Coordinate Output

The internal state representations of the two models are fundamentally different: the lattice model uses a boolean bit-map of occupancy, while the off-lattice model uses complex floating-point numbers with a particle radius $r = 1$ (effective diameter $d = 2$).

To enable a unified analysis pipeline, we implemented a standardisation layer that normalises both outputs into an identical format: a centered $(N, 2)$ array of dimensionless coordinates.

- **Lattice Conversion:** We convert the sparse bit-map into a list of integer coordinates representing the centers of the occupied sites, shifting the seed to the origin $(0, 0)$. We also ensure these are only the occupied cells, as Loh's code includes the sticky sites as well.

- **Off-Lattice and Hybrid Normalisation:** We rescale the continuous coordinates by a factor of 0.5. Since the off-lattice physics uses $d = 2$ and the lattice uses unit spacing $d = 1$, this normalisation ensures that spatial metrics (such as the Radius of Gyration) are directly comparable across models.

This post-processing step allows us to apply the exact same analysis scripts (fractal dimension estimation, radial density functions) to all datasets agnostically, eliminating potential errors arising from manual unit conversion.